

Assignment-4 Use Autoencoder to implement anomaly detection. Build the model by using: a. Import required libraries b. Upload / access the dataset c. Encoder converts it into latent representation d. Decoder networks convert it back to the original input e. Compile the models with Optimizer, Loss, and Evaluation Metrics

Ashish Kumar B511002 BE-A

```
In [1]: #importing libraries and dataset
import numpy as np
import pandas as pd
import tensorflow as tf
import matplotlib.pyplot as plt
from sklearn.metrics import accuracy_score
from tensorflow.keras.optimizers import Adam
from sklearn.preprocessing import MinMaxScaler
from tensorflow.keras import Model, Sequential
from tensorflow.keras.layers import Dense, Dropout
from sklearn.model_selection import train_test_split
from tensorflow.keras.losses import MeanSquaredLogarithmicError

PATH_TO_DATA = 'http://storage.googleapis.com/download.tensorflow.org/data/ecg.csv'
data = pd.read_csv(PATH_TO_DATA, header=None)
data.head()
```

```
Out[1]:
```

	0	1	2	3	4	5	6	7	8	9	...	131	132
0	-0.112522	-2.827204	-3.773897	-4.349751	-4.376041	-3.474986	-2.181408	-1.818286	-1.250522	-0.477492	...	0.792168	0.93354
1	-1.100878	-3.996840	-4.285843	-4.506579	-4.022377	-3.234368	-1.566126	-0.992258	-0.754680	0.042321	...	0.538356	0.65688
2	-0.567088	-2.593450	-3.874230	-4.584095	-4.187449	-3.151462	-1.742940	-1.490659	-1.183580	-0.394229	...	0.886073	0.53145
3	0.490473	-1.914407	-3.616364	-4.318823	-4.268016	-3.881110	-2.993280	-1.671131	-1.333884	-0.965629	...	0.350816	0.49911
4	0.800232	-0.874252	-2.384761	-3.973292	-4.338224	-3.802422	-2.534510	-1.783423	-1.594450	-0.753199	...	1.148884	0.95843

5 rows × 141 columns

```
In [2]: #finding shape of the dataset
data.shape
```

```
Out[2]: (4998, 141)
```

```
In [3]: #splitting training and testing dataset
features = data.drop(140, axis=1)
target = data[140]
x_train, x_test, y_train, y_test = train_test_split(
    features, target, test_size=0.2, stratify=target
)
train_index = y_train[y_train == 1].index
train_data = x_train.loc[train_index]
```

```
In [4]: #scaling the data using MinMaxScaler
min_max_scaler = MinMaxScaler(feature_range=(0, 1))
x_train_scaled = min_max_scaler.fit_transform(train_data.copy())
x_test_scaled = min_max_scaler.transform(x_test.copy())
```

```
In [5]: #creating autoencoder subclass by extending Model class from keras
class AutoEncoder(Model):
    def __init__(self, output_units, ldim=8):
        super().__init__()
        self.encoder = Sequential([
            Dense(64, activation='relu'),
            Dropout(0.1),
            Dense(32, activation='relu'),
            Dropout(0.1),
            Dense(16, activation='relu'),
            Dropout(0.1),
            Dense(ldim, activation='relu')
        ])
        self.decoder = Sequential([
            Dense(16, activation='relu'),
            Dropout(0.1),
            Dense(32, activation='relu'),
            Dropout(0.1),
            Dense(64, activation='relu'),
            Dropout(0.1),
            Dense(output_units, activation='sigmoid')
        ])

    def call(self, inputs):
        encoded = self.encoder(inputs)
```


```
    decoded = self.decoder(encoded)
    return decoded
```

```
In [6]: #model configuration
model = AutoEncoder(output_units=x_train_scaled.shape[1])
model.compile(loss='mse', metrics=['mse'], optimizer='adam')
epochs = 20

history = model.fit(
    x_train_scaled,
    x_train_scaled,
    epochs=epochs,
    batch_size=512,
    validation_data=(x_test_scaled, x_test_scaled)
)
```

```
Epoch 1/20
5/5 ————— 1s 34ms/step - loss: 0.0115 - mse: 0.0258 - val_loss: 0.0135 - val_mse: 0.0313
Epoch 2/20
5/5 ————— 0s 6ms/step - loss: 0.0110 - mse: 0.0247 - val_loss: 0.0129 - val_mse: 0.0300
Epoch 3/20
5/5 ————— 0s 7ms/step - loss: 0.0100 - mse: 0.0223 - val_loss: 0.0126 - val_mse: 0.0290
Epoch 4/20
5/5 ————— 0s 7ms/step - loss: 0.0086 - mse: 0.0191 - val_loss: 0.0129 - val_mse: 0.0294
Epoch 5/20
5/5 ————— 0s 6ms/step - loss: 0.0077 - mse: 0.0168 - val_loss: 0.0121 - val_mse: 0.0278
Epoch 6/20
5/5 ————— 0s 6ms/step - loss: 0.0066 - mse: 0.0146 - val_loss: 0.0120 - val_mse: 0.0276
Epoch 7/20
5/5 ————— 0s 6ms/step - loss: 0.0059 - mse: 0.0130 - val_loss: 0.0117 - val_mse: 0.0268
Epoch 8/20
5/5 ————— 0s 6ms/step - loss: 0.0053 - mse: 0.0118 - val_loss: 0.0113 - val_mse: 0.0260
Epoch 9/20
5/5 ————— 0s 10ms/step - loss: 0.0049 - mse: 0.0110 - val_loss: 0.0108 - val_mse: 0.0250
Epoch 10/20
5/5 ————— 0s 7ms/step - loss: 0.0048 - mse: 0.0107 - val_loss: 0.0104 - val_mse: 0.0241
Epoch 11/20
5/5 ————— 0s 6ms/step - loss: 0.0046 - mse: 0.0102 - val_loss: 0.0100 - val_mse: 0.0232
Epoch 12/20
5/5 ————— 0s 6ms/step - loss: 0.0044 - mse: 0.0098 - val_loss: 0.0098 - val_mse: 0.0228
Epoch 13/20
5/5 ————— 0s 6ms/step - loss: 0.0043 - mse: 0.0096 - val_loss: 0.0096 - val_mse: 0.0225
Epoch 14/20
5/5 ————— 0s 6ms/step - loss: 0.0043 - mse: 0.0095 - val_loss: 0.0096 - val_mse: 0.0224
Epoch 15/20
5/5 ————— 0s 6ms/step - loss: 0.0042 - mse: 0.0095 - val_loss: 0.0096 - val_mse: 0.0223
Epoch 16/20
5/5 ————— 0s 6ms/step - loss: 0.0041 - mse: 0.0092 - val_loss: 0.0095 - val_mse: 0.0222
Epoch 17/20
5/5 ————— 0s 6ms/step - loss: 0.0040 - mse: 0.0090 - val_loss: 0.0094 - val_mse: 0.0221
Epoch 18/20
5/5 ————— 0s 6ms/step - loss: 0.0039 - mse: 0.0088 - val_loss: 0.0094 - val_mse: 0.0219
Epoch 19/20
5/5 ————— 0s 6ms/step - loss: 0.0040 - mse: 0.0089 - val_loss: 0.0093 - val_mse: 0.0217
Epoch 20/20
5/5 ————— 0s 6ms/step - loss: 0.0037 - mse: 0.0084 - val_loss: 0.0092 - val_mse: 0.0215
```

```
In [7]: plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.xlabel('Epochs')
plt.ylabel('MSLE Loss')
plt.legend(['loss', 'val_loss'])
plt.show()
```



```
In [8]: #finding threshold for anomaly and doing predictions
def find_threshold(model, x_train_scaled):
    reconstructions = model.predict(x_train_scaled)
    reconstruction_errors = tf.keras.losses.msle(reconstructions, x_train_scaled)
    threshold = np.mean(reconstruction_errors.numpy()) \
        + np.std(reconstruction_errors.numpy())
    return threshold

def get_predictions(model, x_test_scaled, threshold):
    predictions = model.predict(x_test_scaled)
    errors = tf.keras.losses.msle(predictions, x_test_scaled)
    anomaly_mask = pd.Series(errors) > threshold
    preds = anomaly_mask.map(lambda x: 0.0 if x == True else 1.0)
    return preds

threshold = find_threshold(model, x_train_scaled)
print(f"Threshold: {threshold}")
```

73/73 ————— 0s 1ms/step
Threshold: 0.008325223361028106

```
In [9]: #getting accuracy score
predictions = get_predictions(model, x_test_scaled, threshold)
accuracy_score(predictions, y_test)
```

32/32 ————— 0s 613us/step

Out[9]: 0.947